School of Electronic Engineering

CB54: Machine Learning Algorithms for EM Wave Scattering Problems

Appendix E: Testing & Results

Anthony James McElwee

ID Number: 20211330

August 2023

MEng in Electronic and Computer Engineering

Supervised by Dr Conor Brennan

Contents

Introduction

All statistical tests were performed using the free and open-source statistical analysis programme JASP [1]. The JASP programmes and the associated CSV datasets are available on GitHub. A simple Bayesian alternative to the frequentist approach also resides within the JASP files.

Description of Training Dataset

The generation of seeded samples was conducted in batches of one thousand samples per folder due to the input configuration used in Python causing memory exhaustion. Five thousand samples were generated. All folders were scanned using Auslogics Duplicate File Finder 10, on the PNG files, to search for duplicate geometry samples. These duplicates and their associated NumPy files were moved to a separate duplicate folder and excluded from all further experimental activity. Earlier runs of the experiment had used an even more minimal geometric scenario generator where the main scatter was anchored at the origin. After generating 49000 samples, and extremely late in the project timeline, it was realised by the student that roughly 84% of the samples were duplicates. The training, validation and testing sets were all totally overlapping and the model was immediately overfitting on every run. After correcting for this issue, the student found that between 4000 and 5000 samples were enough to develop the model before the training loss curve started to indicate possible overfitting issues. The training set was split at 80% from the entire folder with the remaining 20% used for testing. A further split of 20% from the training set was used for validation at the end of each epoch leaving 640 samples for training in each session. Due to the removal of duplicates some folders had less than 640 training samples.

E1 Model Training Commentary

The final total model fitting time for the E1 component of the Prescient2DL deep learning model was 1117.3092517852783 seconds, which is roughly twenty minutes. This does not include the initial creation of the training/validation/testing splits which take time to transform from the sample data to correct tensor format. The screen command line printout of the final two training data batches is available on GitHub for reference.

E1 Field Model Training Loss Curves

It should be noted that GitHub hosts all of the loss curves plotted at the end of each epoch. Tensorboard was also used to track the weight updates at each layer, however, analysis and improvement on weight behaviours will remain in the domain of future work due to time limitations.


Figure 1Final Loss Curve Plot. Note the axis scale is logarithmic.

The loss curve is extremely severe when plotted on the normal axis and corresponds with loss curve plots found in the paper [2] that deals with physics deep learning applications to a laminar flow problem in the domain of fluid dynamics. Both training and validation curves track each other well In the final training session, the training curve starts to oscillate compared to the validation curve indicating possible overfitting. Although the loss curve generally follows the desired shape of a loss curve graph, albeit on a log scale, the loading spikes indicate a number of possible issues in the model design. These issues are discussed later in the appendix.

E1 Field Model Training Test Scores

For the final training dataset, the following scores were achieved:

Metric	#
Training Mean Squared Error Loss	2.3042e-05
Validation Mean Squared Error Loss	2.5000e-05
Final Learning Rate	1.0000e-22

Key Metric: For the final training dataset, the mean squared error test score after training was 2.628892798384186e-05. Since no cross-contamination of samples exist in each split, the testing error here is the correct indicator of model performance from the perspective of the deep learning model compared to the final solved solution provided as the prediction target. The final learning rate frequently was reduced to this order of magnitude throughout each training dataset run and this is reflected in the loss curves.

E1 Field Model Training Visual Prediction Results Absolute Component

The student collected prediction plots of a sample test case at the end of each training set loading stage to illustrate the improvement in the prediction power of the model as it was being trained. Each dataset contained roughly 640 samples of training data, depending on duplicates, after the testing and validation splits were removed. The plots are shown below in chronological order and can be found on GitHub. They illustrate the input geometry contrast values and the absolute field assembled form the two predicted real and imaginary component fields. The rapid convergence to low mean squared error loss is clearly captured in the difference plot on the bottom left corner, however, the source of truth and predicted fields really only become comparable as the model starts to overfit in the final two plots.


Figure 2End of first training session for E1.

Figure 3. End of second training session for E1.

Figure 4. End of third training session. For E1.

Figure 5. End of fourth training session for E1.

Figure 6. End of fifth training session for E1.

E2 Model Training Commentary

The final total model fitting time for the E2 component of the Prescient2DL deep learning model was 1183.819656610489 seconds, which is roughly twenty minutes. This is almost the same as the E1 component training time and this is to be expected as the same architecture was used in both cases. This does not include the initial creation of the training/validation/testing splits which take time to process from the sample data to correct tensor format. The screen command line printout of the final two training data batches is available on GitHub for reference.

E2 Field Model Training Loss Curves

It should be noted that GitHub hosts all of the loss curves plotted at the end of each epoch. Tensorboard was also used to track the weight updates at each layer, however, analysis and improvement on weight behaviours will remain in the domain of future work due to time limitations.

Figure 7. Final Loss Curve Plot E2. Note the axis scale is logarithmic.

E2 Field Model Training Test Scores

For the final training dataset, the following scores were achieved:

Metric	#
Training Mean Squared Error Loss	3.7006e-05
Validation Mean Squared Error Loss	3.7967e-05
Final Learning Rate	1.0000e-22

Key Metric: For the final training dataset, the mean squared error test score after training was 4.160570097155869e-05. Since no cross-contamination of samples exist in each split, the testing error here is the correct indicator of model performance from the perspective of the deep learning model compared to the final solved solution provided as the prediction target. This is notably higher than the E1 score, however, the complexity of the y-dimension is more apparent since the incident dipole wave has greater magnitude of change compared to the x-dimension co-ordinate. The final learning rate frequently was reduced to this order of magnitude throughout each training dataset run and this is reflected in the loss curves.

Figure 8. End of first training session for E2.


Figure 9. End of second training session for E2.

Figure 10. End of third training session for E2.


Figure 11. End of fourth training session for E2.

Figure 12. End of fifth training session for E2.

E-9

Collected comments on the E1 & E2 Training Process

It appears that both models are not learning in a manner usually desired in deep learning development processes. The student will now provide some commentary on this topic.

The search space of possible geometric configurations for this problem can be approximated with an upper bound on the number of cells in the grid multiplied by the possible starting positions for the smaller scatterer in the scenario. This can be roughly estimate to be (128^2 * 1600) which gives 26 million possibilities. Five-thousand geometric samples is roughly 0.02% coverage of this space. The means squared error results indicate that the deep learning models are successfully capturing a compressed version of the fields in their weights but only up to a certain degree of accuracy. Based on the oscillating training curve in the final run of the training, it is unlikely that there is insufficient training data. Indeed, more training data could be generated and up to 49000 data samples had been generated in earlier conceptions of the project development. The use of training augmentations is also eliminated since medical applications require pre-designated incident directions. Re-orientating the incident fields will not increase model generalisability that is useful to the long term goal of the research. With the exception of horizontal mirroring, data augmentation would shrink the possible permutations in the scene configuration space and increase the probability of duplicates between the training/test/validation sets.

The student finds it much more likely, based on the literature review, that the U-net architecture in its current form is sub-optimal or inappropriate for the problem at hand. The student has reflected on the use of meta-architecture in the project research log and this would be a route that research could follow off the back of the findings made in this project.

At time of writing, the student has begun to interpret the general application of deep learning to the forward problem using U-Net as an implicit recasting of the problem to an inverse one where the inputs are basically noisy inputs that need to be recovered backwards to their scattered fields. This had been the students long term outlook when conducting the literature review, except using different prior information gleamed from Von Neumann integral expansions in the higher frequency approximations to the wave scattering configurations and then applying denoising to complete the scene. There is a conceptual connection in this idea to the GANs approach documented in [3].

Although the student cannot rule out the possibility that the model has received inadequate training time, the rapid convergence to a similar order of magnitude of loss through all training sessions implies that this is not a training time duration issue. The student also believes increasing the training data size to cover more of the possible permutations would not scale well when higher-contrast problems are tackled since this would dramatically increase the amount of training data required to achieve results, in direct contradiction to developing the model in the first instance. It is a key requirement that only a small sample set is required to train and test the model, as this is what would be available to researchers at the low-frequency high-contrast end of the problem spectrum.

While the predicted target data was pre-processed and then post-processed to bring the values within a smaller range suitable for deep learning, this was not done for the input data. For this low-contrast problem where the incident waves are already tightly bounded in value, it is unlikely that this is causing a problem. However, the student notes that in the high-contrast scenarios such processing would also need to be carried out on the input data in order for the model weights to adjust quickly and avoid losing permutations to bias adjustment. As already commented, the model incorporates a bias term in each convolution layer but this bias term still needs to learn so processing the input data further may help this delivery. The student has included many batch normalisation layers in the architecture in a bid to alleviate this issue.

The student has adapted a learning rate that reduces over time as the loss curve saturates. The use of an increased batch size compared to [4] up to where the computer memory would allow should also have aided the improved training performance. The model also had two extra final linear layers that aimed to provide a blurring effect on the
E-10

arrays, given the smoothness assumption for the scattered fields. The student also included a Dropout layer need the lower end of the decoding side of the model. The student notes that the use of batch layer normalisation may help the model to train more smoothly, however, implicit regularisation may also be arising from these layers and the model may be struggling to reduce the loss values with higher accuracy because of their inclusion at so many levels. The student also did not use advanced hyper-parameter tuning or meta-search libraries such as AutoKeras. In general, the student would suggest experimenting with regularisation on the activity of some of the convolution layers should further research be conducted on this specific model architecture.

Although not formally reported, the student tried to use Xavier weight initialisation as suggested in [4]. This did not positively impact the learning curves. The student used Adam optimiser based on the literature review.

Finally, the student wants to raise the possibility that the application of deep learning to this problem may not be suitable beyond the emulation case. This project has failed to find evidence that infusing deep learning models into conventional Krylov based solvers can improve their convergence properties. While the model has been shown to provide a decent estimation of the target fields, substantial benefits to existing methodologies can only be claimed if either the initial error arising from using the deep model lower the iteration count of the Krylov solver or help the Krylov solver to converge at a faster rate to a solution that meets the error criterion of the simulation. Evidence for either of these goals was not found in the experiments conducted in this project. This is evidenced in the next section.

Descriptive Statistics of Testing Datasets

The following table gives the descriptive statistics for each statistical test set used to evaluate the impact of Prescient2DL on SolverEMF2. Each set consisted of 100 original samples solved using the naïve initial guess of the incident wave as the scattered field. None of these samples appeared in the training/validation or test sets used during model development. After training the models for predicting the two scattered fields, a second run of SolverEMF2 was used on the same original samples, allowing for direct comparison across duration of calculation, iteration count and initial error between the original sample information ("_o") and the model-assisted sample information ("_m").

| Metric	N	Mean	SD	SE	Coefficient of variation
DS1					
Duration_o	100	1.106213	0.055213	0.005521	0.0499113
Duration_m	100	1.010293	0.093377	0.009338	0.0924252
Iteration_Count_o	100	22.57	0.655282	0.065528	0.0290333
Iteration_Count_m	100	22.05	0.479373	0.047937	0.0217402
Error_Initial_o	100	0.004857	0.002712	0.000271	0.5583008
Error_Initial_m	100	0.001102	0.000492	4.92E-05	0.4466384
DS2					
Duration_o	100	0.772	0.132	0.013	0.171
Duration_m	100	0.72	0.071	0.007	0.099
Iteration_Count_o	100	19.57	0.573	0.057	0.029
Iteration_Count_m	100	19.35	0.52	0.052	0.027
Error_Initial_o	100	0.003	0.002	1.963×10^{-4}	0.619
Error_Initial_m	100	8.014×10^{-4}	3.702×10^{-4}	3.702×10^{-5}	0.462
DS3					
Duration_o	100	2.218	0.198	0.02	0.089
Duration_m	100	2.308	0.246	0.025	0.106
Iteration_Count_o	100	56.65	1.048	0.105	0.019
Iteration_Count_m	100	56.58	0.955	0.096	0.017
Error_Initial_o	100	0.03	0.019	0.002	0.611
Error_Initial_m	100	0.02	0.011	0.001	0.533

Paired t-Tests of Testing Datasets

Between the descriptive statistics of the testing datasets and the t-tests conducted on the variables, the impact of Prescient2DL on Solver EMF2 can be established. In the attached appendix Project Plan Proposal, the Primary research hypothesis were tested. The Primary Research question was orientated around using Prescient2DL to improve the Krylov Iterative solver at the heart of SolverEMF2. DS1, DS2 and DS3 testing datasets had 100 samples each which gives enough degrees of freedom to ensure a large enough sample size to compare the original solution information to the information gathered during the model assisted solution. Please review the Project Plan Proposal for the details of the three hypothesis: "Primary Research Test 01 – Initial Solution Conveyance t-Test"; "Primary Research Test 02 –Solution Convergence t-Test"; "Primary Research Test 03 – Solution Conveyance t-Test".

DS1 Results & Analysis

DS1: Paired Samples T-Test								
Measure 1	Measure 2	t	df	p	Mean Difference	SE Difference	Cohen's d	SE Cohen's d
Duration_o	Duration_m	8.9132305	99	< .001	0.0959198	0.0107615	0.8913231	0.1656468
Iteration_Count_o	Iteration_Count_m	7.2478005	99	< .001	0.52	0.0717459	0.7247801	0.1394524
Error_Initial_o	Error_Initial_m	16.5942404	99	< .001	0.0037544	0.0002263	1.659424	0.0587952

In the results for DS1, all p-values were deemed statistically significant. However, looking beyond the t-test results and comparing the mean differences for the Duration and Iteration Count variables, the impact the model has on Solution Convergence and Solution Conveyance is inconsequential to the solution progress of the Krylov Iterative Solver. On average, the same number of iterations are required to meet the error criterion of the solver. This is reflected in the Raincloud plots for the Duration and Iteration count variables.

Figure 13. Raincloud plots illustrating the similiarity between the original and model-assisted samples. Although statistically different, the lowering of the variables by the model is immaterial to the solver.

Figure 14. Raincloud plot of initial error comparison. While the initial error in the model-assisted approach is lower and more consistant then the naive approach, it is not low enough to impact SolverEMF2.

The Initial Solution Conveyance based on a comparison of the means of the Error Initial variable is statistically significant with a seemingly impressive mean difference of 0.0038. This is roughly 18% of the original initial error in the naïve approach. As demonstrated, this is enough to give a close visual similarity that would be useful in the early stages of a design methodology where error criteria are low. Although the statistical test rejects the null, indicating a clear performance difference exists between the two approaches, the student concludes that the impact of Prescient2DL is immaterial to achieving the error criterion set out at the initialization stage of the solver.

DS2 Results & Analysis

The aim of the DS2 dataset was to establish how Prescient2DL would handle geometric configurations beyond those it was exposed to in the training stage of development. By setting the smaller scatterer to the same zero contrast value as the background embedding, a deformation or hole was created in the main scatterer geometry.

DS2: Paired Samples T-Test								
Measure 1	Measure 2	t	df	p	Mean Difference	SE Difference	Cohen's d	SE Cohen's d
Duration_o	Duration_m	3.394	99	< .001	0.052	0.015	0.339	0.15
Iteration_Count_o	Iteration_Count_m	4.2	99	< .001	0.22	0.052	0.42	0.1
Error_Initial_o	Error_Initial_m	14.493	99	< .001	0.002	1.634×10^{-4}	1.449	0.061

Again, all p-values indicate a rejection of the null hypothesis. Prescient2Dl is able to operate with this mild generalization. All three metrics have reported a decrease in the mean difference between naïve and model-informed runs, however, the decrease for the Error Initial metric is mild with a score of roughly 27% compared to the initial error offered by the original, naïve approach.

DS3 Results & Analysis

The aim of the DS3 was to push Prescient2DL into predicting scenarios where the contrast values populating the scene were significantly different to those it was exposed to during training. In this case, both the small and larger scatterers took the same higher contrast value.

DS3: Paired Samples T-Test								
Measure 1	Measure 2	t	df	p	Mean Difference	SE Difference	Cohen's d	SE Cohen's d
Duration_o	Duration_m	-3.3	99	0.001	-0.09	0.027	-0.33	0.125
Iteration_Count_o	Iteration_Count_m	0.572	99	0.569	0.07	0.122	0.057	0.122
Error_Initial_o	Error_Initial_m	12.829	99	< .001	0.01	8.095×10^{-4}	1.283	0.017

The Duration metric indicates that the model actually has a detrimental impact on SolverEMF2. The evidence supplied to the t-test results in an inability to reject the null hypothesis for Iteration Count. As already commented in DS1 and DS2, neither of these mean differences are material to the Krylov solver in the grand scheme of the workflow. The success measured by the Error Initial variable has also deteriorated but still reports roughly 66% of

the original, naïve approach. Even though Prescient2DL was not exposed to single scatterer geometries with such a high population of higher contrast values, it was still able to lower the initial error. A sample prediction outcome is presented in the figure below. The titles indicate the figure content. It is evident that the visual resemblence between the predicted fields and the actual fields has broken down, even if the residual error remains lower than that of a naïve guess.

Figure 15. Comparison between the actual scattered field on the left and the predicted field on the right for E1.

Figure 16. Comparison between the actual scattered field on the left and the predicted field on the right for E2.

Impact Demands

If the current tests and results show low means squared error but immaterial impact upon SolverEMF2, how would results that actually achieve Method of Moments solution acceleration manifest in this scenario? Using a DS1 sample case, the student prints out the first seven entries from the information gathered during a naïve and model-assisted run in the table below.

Sample 5101	Naïve Original Run		Prescient2DL Assisted Run	
Iteration	Residual Error	Duration	Residual Error	Duration
0	0.0020748100	0.0102043000	0.0005120250	0.0085418200
1	0.0002348990	0.0905674000	0.0000702384	0.0611167000
2	0.0000368364	0.1306520000	0.0000158675	0.0931430000
3	0.0000102315	0.1706970000	0.0000042492	0.1249560000
4	0.0000010167	0.2137400000	0.0000005663	0.1580040000
5	0.0000002997	0.2488170000	0.0000000635	0.1890330000
6	0.0000000401	0.2823450000	0.0000000171	0.2185600000

Table 1. Iterative solver information for sample 5101 using DS1 scene parameters. Only first 7 entries are displayed.

In order for Prescient2DL to impact the Method of Moments solver by accelerating the solution process, the initial residual error would need to be lowered below at least these severn error levels reported in the residual error column for the Naïve Original Run. Due to dependance of convergence rates for Krylov solvers on the conditioning of the matrices and the eigenvalue properties of the matrices, deep learning models would need to be achieving initial residual errors of 10^{-8} or lower in this toy scenario to lower the final iteration count by even 25% of the naïve solution iteration count.

Conclusions

Although Prescient2DL is statistically achieving differences in metric performance compared to naïve approaches to the initial guess in SolverEMF2, there appears to be nothing gained from using the model at all with regard to accelerating the Method of Moments solver. Plots of convergence and conveyance, as postulated in the Project Plan Proposal, would all be identical between original and model-assisted runs. Before further experimentation with generalizability and extrapolation is to be carried out, it would be the student's opinion that lowering the initial error so that Prescient2DL impacts the SoverEMF2 in terms of convergence would be made the priority. This project has presented evidence and test results that show research in this area remains at the model development stage. If the creation of models that bring the initial residual error to orders of magnitude lower than currently reported is deemed non-viable, after more extensive experimentation, then the attempted application of deep learning in electromagnetic scattering forward problems could be regarded as frivolous.

Bibliography

[1] JASP Team, "JASP (Version 0.17.3)." 2023. [Online]. Available: https://jasp-stats.org/

[2] T.-T.-H. Le, H. Kang, and H. Kim, "Towards Incompressible Laminar Flow Estimation Based on Interpolated Feature Generation and Deep Learning," *Sustainability*, vol. 14, no. 19, Art. no. 19, Jan. 2022, doi: 10.3390/su141911996.

[3] Z. Ma, K. Xu, R. Song, C.-F. Wang, and X. Chen, "Learning-Based Fast Electromagnetic Scattering Solver Through Generative Adversarial Network," *IEEE Trans. Antennas Propag.*, vol. 69, no. 4, pp. 2194–2208, Apr. 2021, doi: 10.1109/TAP.2020.3026447.

[4] Q. Ren, Y. Wang, Y. Li, and S. Qi, *Sophisticated Electromagnetic Forward Scattering Solver via Deep Learning*. Singapore: Springer, 2022. doi: 10.1007/978-981-16-6261-4.